
SPIKING NEURAL NETWORK DRONES

PROJECT DESIGN REPORT

December 6, 2023

University of Tennessee, Knoxville

COSC 401 Senior Design

Team 1

Tyler B. Nitzsche, Jihun Kim, Seoyoung An, Jonathan E. Skeen

Advisor

Dr. Catherine D. Schuman

Executive Summary

The purpose of this project is to train and simulate a Crazyflie drone using a spiking neural network so that it can operate in a real-world environment by itself. Our goal is to demonstrate a drone to take off, hover, and land by itself.

Simulation is the first step to a spiking neural network drone. Spiking neural networks are inspired by biological brains and take information from a given state to make computations based on that given data. By using a genetic algorithm, we will build a spiking neural network that is optimized to get the optimal output. We will use EONS, Evolutionary Optimization for Neuromorphic Systems developed by TENNLab, for model training. EONS utilizes evolutionary optimization for the applications of spiking neural networks, and we will use it to train our drone in a simulation environment to check its performance.

To test it, we will use a virtual simulation environment. We have three options: the pre-existing environment from last year's team, a simulator and an environment from Dr. Simon D. Levy from Washington and Lee University, and to develop our own simulation environment if neither of these work.

After testing it in a virtual environment, we will start to interact with the hardware of the drone and test it with a physical drone. Due to Federal Aviation Regulations, we will notify McGhee Tyson Airport and the University of Tennessee Police Department at least a week before to request that we can fly the drone.

Each step of observation with the sensor, training with networks, and making decisions based on the networks will be repeated until we reach the goal: a drone that can successfully take off, hover, and land by itself without any problem.

Table of Contents

Executive Summary	1
Problem Definition & Background	4
The Requirement Specification	6
Technical Approach	8
Design Concepts, Evaluation & Selection	9
Deliverables	10
Project Management	11
Roles	11
Milestones	12
Budget	13
Reference	14

List of Figures

1	Crazyflie 2.1 Drone [1]	4
2	Comparison between traditional Von Neumann and Neuromorphic architectures [2]	5
3	Overall Project Process	6
4	Overview of EONS Algorithm [3]	7
5	Part of the Gantt Table of the Project	13

6	Project Quad Chart	15
---	------------------------------	----

List of Tables

1	Milestone in Each Sprint	16
---	------------------------------------	----

PROBLEM DEFINITION & BACKGROUND

Autonomous drones have been a goal the artificial intelligence community has been working towards for decades at this point. Whether this idea is marvelous or terrifying, with the recent spark in interest for artificial intelligence, it seems as if the advancement of this technology is inevitable. Thus, our team has decided to approach this problem with neuromorphic computing. More specifically, we will train a spiking neural network to operate a small drone, performing a variety of simple tasks.



Figure 1: Crazyflie 2.1 Drone [1]

When approaching this problem today, one of the most popular approaches is deep learning – a subset of traditional neural networks. Traditional neural networks, as a whole, have seen an enormous jump in popularity over recent years, mainly due to their remarkable performance improvements. When comparing machine learning approaches for problem-solving, often deep learning is the path that many developers will pursue. While there are great results that people have had with deep learning, the downsides to these networks can outweigh the benefits.

While there is a continually growing list of concerns about using deep networks (such as the enormous amounts of data required to train these networks), our team’s main concerns with using deep learning are its size and efficiency. Our network will need to operate within the drone’s hardware. Thus, the network’s storage and computational resources are severely limited, making a deep-learning approach unsuitable.

Enter, the spiking neural network (SNN). These networks operate differently when compared to traditional neural networks. Both types share the idea of neurons and synapses/edges, but spiking neural networks attempt to do more closely replicate biology by incorporating time as a component. In addition, SNNs often have fewer neurons and edges, helping with the storage concerns, furthermore, they have relatively lower power consumption (due to neurons being idle most of the time), helping with the computation concerns. These networks do, however, require the use of a specialized piece of hardware called a neuroprocessor, but since we already have this hardware, this isn’t much of a concern.

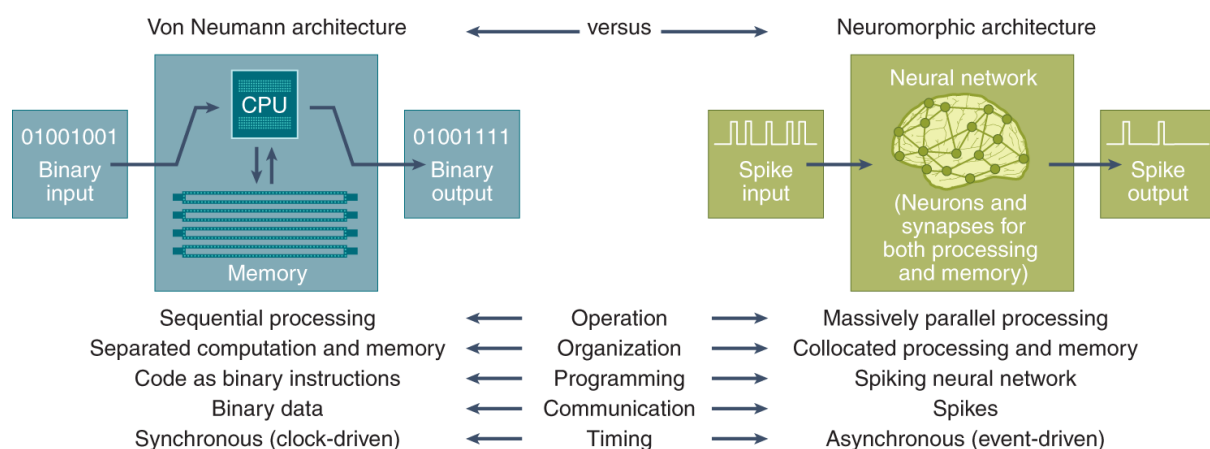


Figure 2: Comparison between traditional Von Neumann and Neuromorphic architectures [2]

The last consideration for this problem is determining the network structure. Even with limited inputs and outputs (less than 20 in total), there could be thousands or millions of possible networks. Using SNNs makes these issues worse since these networks are usually not grid-like nor fully connected.

Yet another problem to consider for this task is the lack of knowledge regarding neural networks in general. Neural networks lack a key component essential for many problem-solving tasks: explainability. As of now, it is very difficult to accurately determine exactly how network structure will affect output and the network also cannot provide an explanation. Given all of this information, it becomes apparent that manually solving this problem is not optimal. Fortunately, Dr. Schuman and TENNLab have developed an algorithm that helps with creating model structures, tailored specifically to SNNs: Evolutionary Optimization for Neuromorphic Systems (EONS) [3]. Another algorithm attempting to replicate biology, EONS combines randomness and performance to enhance and modify spiking neural networks. Thus, this is our key to creating and training SNNs [3].

THE REQUIREMENT SPECIFICATION




Figure 3: Overall Project Process

Our requirement specifications can be broken up into two main parts: Simulation and real-world specifications. Because our physical hardware is expensive and can pose a safety risk, the majority of our project will pertain to a simulation environment. This environment must contain three things to successfully fulfill our needs: the ability to simulate the drone within the environment, all relevant and useful data required by our SNN, and a high similarity to the real-world to ease the transition process from the simulation to the physical environment. Thus, the first step in our process is to find the best simulation environment amongst the three options provided to us by Dr. Schuman. After an environment is chosen, our team must

incorporate our drone into the environment. This process includes ensuring proper flight operation as well as proper sensor readings, as described in Fig 3.

After the preparation tasks are finished, we will begin developing various fitness functions for our networks. Maximizing these functions is the ultimate goal of SNNs and is our way of telling a network how it's performing. Correctly and accurately defining these functions is a notoriously difficult problem and will likely take up a sizeable chunk of our time. Since we are ambitious, we plan to define multiple fitness functions, which each correlate to a different task.

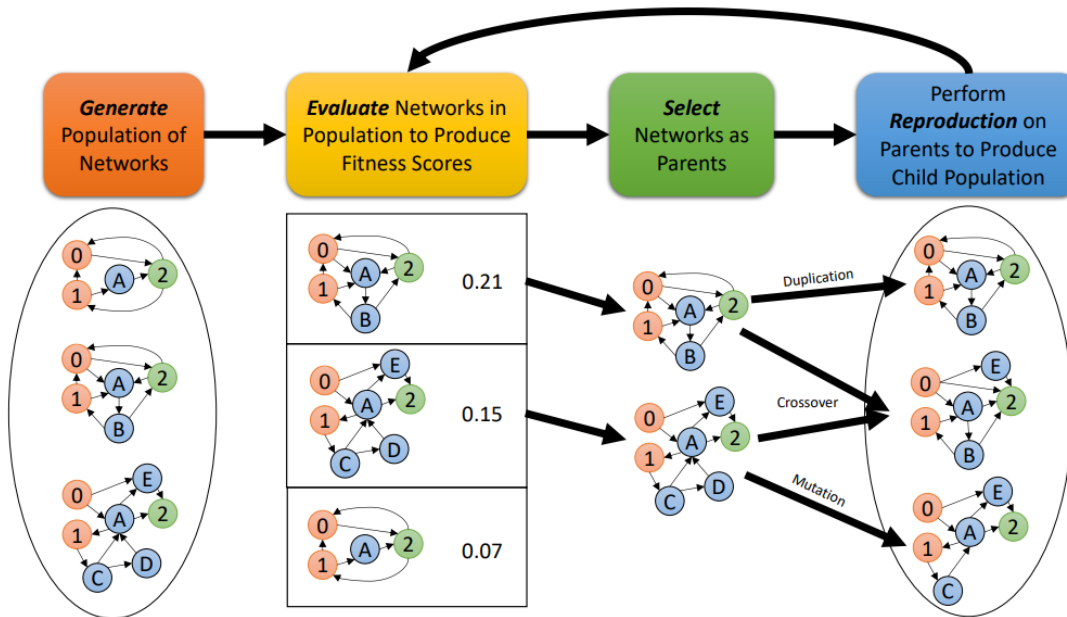


Figure 4: Overview of EONS Algorithm [3]

Once we have fine-tuned our fitness functions, we will implement the EONS algorithm into the simulation and begin training our networks [3]. Figure 4 shows each step of EONS [3]. Once again, we anticipate that this phase will also take up a large portion of our time [3]. If there is success in the simulation environment, we will begin the transition into real-world operation.

Transitioning into the physical environment will require a few different steps, with the first

involving connecting to the drone API. This step could be completed in the simulation phase (when implementing the drone into the simulation environment), but if not, it will be the first task after finishing the simulation phase. From the simulation, we expect to get a few different network structures. We will load each of these networks onto the drone and measure how they perform in the real world. Just as with the simulation, we will have multiple fitness functions, ranging from very simple, to slightly complex, but the tasks will remain basic.

TECHNICAL APPROACH

The project consists of three different components, all of which are unique and warrant their own approaches. The first component is the simulation environment, which is responsible for hosting the training of the drone's algorithm. Then there is the evaluation function, which is expected to comprise the bulk of the development time for the project and controls what the machine learner values in training. Finally, there are hardware considerations and adjustments to ensure that the physical drone matches its virtual counterpart.

For the simulation environment, there are three predominant options that have been considered. The first option is to develop an entirely novel simulation environment that is tailored to the testing of our machine-learning model. The second is to use a simulation environment developed by Simon Levy specifically for the Crazyflie drone [4]. Third is the option which is the one which is currently being most heavily considered, and that is using the simulation environment developed by the prior senior design class group who worked on this project. This third option is the primary consideration because it combines many of the benefits of the prior two options since it is both designed for this specific problem and does not require much more development time.

The evaluation function is the main way the software developed in this project will interact with the drone. It consists of a function that uses the data available to the drone to mathematically determine if the drone is currently flying well. The development of an evaluation function is a complex series of decisions in which the data generated by the model's actions are evaluated to determine how 'well' the model should be considered to be performing. These adjustments allow for the model to assess itself and improve in accordance with the function.

The hardware considerations and decisions will begin to be made later in the project after

some preliminary models have been trained and evaluated, and consist mostly of determining which ways the physical drone needs to be changed in order to align with the simulated version, or if the simulation needs to be altered to match the physical version. Furthermore, there are decisions to be made on which additional components need to be added to the physical drone in order to allow it to have added functionalities, such as reading more data points.

These three design components combine to create an autonomously operating drone, the simulation allows the model to train, the evaluation function tells the simulation how successful the training is, and the hardware is the application of that simulated model. The end result should be a Crazyflie drone able to hover stably a few feet above the ground.

DESIGN CONCEPTS, EVALUATION & SELECTION

For the simulator considerations usage of the prior senior design team's simulator is the main contender. The alternatives are using Simon Levy's simulator or creating a new simulator [5]. There are three main factors to choosing the simulator, that being how well suited it is to the problem, ease of use, and development time. Currently the choice leans towards the previous senior design team's simulator as it was developed specifically for this project, likely suits the project's needs very well, and will theoretically require little to no development time devoted to tweaking the simulator. Much of the same is true with Simon Levy's simulator, however as it was not designed for this project specifically, it is likely it is not as well suited to the project, and therefore the senior design simulator is preferable, however further testing of the two simulators is needed to determine which is better. The third option of development of a new simulator is very unlikely to be employed, while it is possible that a simulator that is more effective and easier to use would be created, such an effort would likely consume all the development time and leave the primary objective of training a drone to hover to be uncompleted [5].

The decisions in the evaluation function cannot yet be properly listed and evaluated in this preliminary stage. Determining which decisions exist to consider and weighing their values will comprise a large portion of the problem space for this project, if a complete description of these decisions could be presented and elaborated upon here, the project would likely be completed. However potential features of interest can be determined, such as the height of the drone, acceleration of the drone, and gyroscope readings. All of these

and more will likely be relevant features, but without extensive testing, it is impossible to be certain what features the drone's model will use.

The decisions associated with hardware will, similarly to the evaluation function, need to be determined by extensive testing, some hardware adjustments will become necessary to accommodate new sensors for the drone to utilize in its data gathering so it can accurately reflect the sensing capabilities of its virtual twin. Other hardware adjustments will need to be made in order to keep it more in line with this virtual twin and make sure it is performing similarly enough to the simulation. Both of these decision points require a somewhat functioning model and a preliminary evaluation function to begin being considered, therefore it is impossible to say with certainty what these decisions look like at this preliminary stage.

The broadest possible hardware consideration that is being considered is choice of drone. Currently the Crazyflie drone is being considered for its small size and already established code base, but should it prove uncooperative other autonomous drone options exist, such as the flapper drone, but the flapper drone has the issue of being much more experimental, as well as expensive, furthermore it has a much more unfamiliar code base, that our researched simulators cannot accommodate. If other drones were implemented it would likely restructure the entirety of the project, and as such should be only considered if usage of the Crazyflie proves to be completely impractical.

DELIVERABLES

The primary deliverable of this project is providing the customer with a drone that is controlled autonomously using spiking neural networks. The initial goals of the drone are to achieve stable drone operation in a simulated environment and achieve stable drone operation in a real-world environment. The stretch goals of the drone are to make the drone safely follow a leader drone and once this is working, we want to apply swarming together. We plan to implement the goals in small drones, and once we are successful in implementing them, we plan to ensure that the features work well in full-size drones. The drone will be delivered with the following components:

1. Crazyflie drone
2. Evolutionary Optimization for Neuromorphic Systems(EONS) for model building and training

-
3. Simulation environment for training and experimenting with the drone and the network
 4. Neuromorphic Hardware to run the spiking neural network controlling the drone

Our Team will choose a small Crazyflie drone which will be the best fit for the project and we will apply our software to the drone to see how it works well. Considering the size and specifications of the drone, we will deliver software and hardware optimized for the drone. Software development will be based on the existing network which needs some control refining. In addition, we will apply EONS to train networks for an application. To test and develop our software, we will use the existing physics simulator which is for takeoff and landing [3].

After we finish our software development, we will adjust the spiking neural network model into the hardware and test our drone. We will ensure the performance, stability, and flight skill of the drone. All the deliverables will be provided to the customer before the end of the spring 2024 semester (May 2024.)

PROJECT MANAGEMENT

Roles

Our project, which consists of four people, aims to have each member choose a role they are confident in, faithfully carry out their duties, and coordinate opinions with each other to create optimal work.

As software developers, all team members will understand how the previous network developed and then work to improve it. We will focus on the process of training a Spiking Neural Network using EONS, and the implemented Spiking Neural Network will be tested using a physical simulator [3] [4] [5].

As a team leader, Tyler Nitzsche will monitor the overall quality of the project. His experience in conducting projects using Spiking Neural Network could be useful in our project. Based on his experience and knowledge, we will draw a blueprint for building a Spiking Neural Network model.

As a project organizer, Jihun Kim will plan the overall schedule of the project and rearrange it considering the circumstances of each group member. To be able to complete a large amount

of work over the next semester, we need to plan the project well.

As hardware developers, Seoyoung An and Jonathan Skeen will apply the software we developed to hardware so that drones can utilize it. To achieve this, they will be responsible for designing and designing the drone's hardware and choosing what hardware to use. This task will be treated as one of the important tasks of the project as it deals with the practical use of the software.

This project covers a lot of tasks, and given the different experiences of each group member, we aim to create an environment where each member not only plays his or her role but also supports each other during the work. Each member will also be a reviewer or tester, finding problems and improvements in each other's work. This work can help prevent bigger problems and create better software or hardware.

Milestones

To ensure deadlines and manage schedules from the Fall 2023 semester to the Spring 2024 semester, we divided the work into five sprints as Table 1 shows. There may be changes in the order of work, additions, or exclusions of tasks in the future.

The Gantt chart in Figure 5 shows the overall schedule of the project in a table.

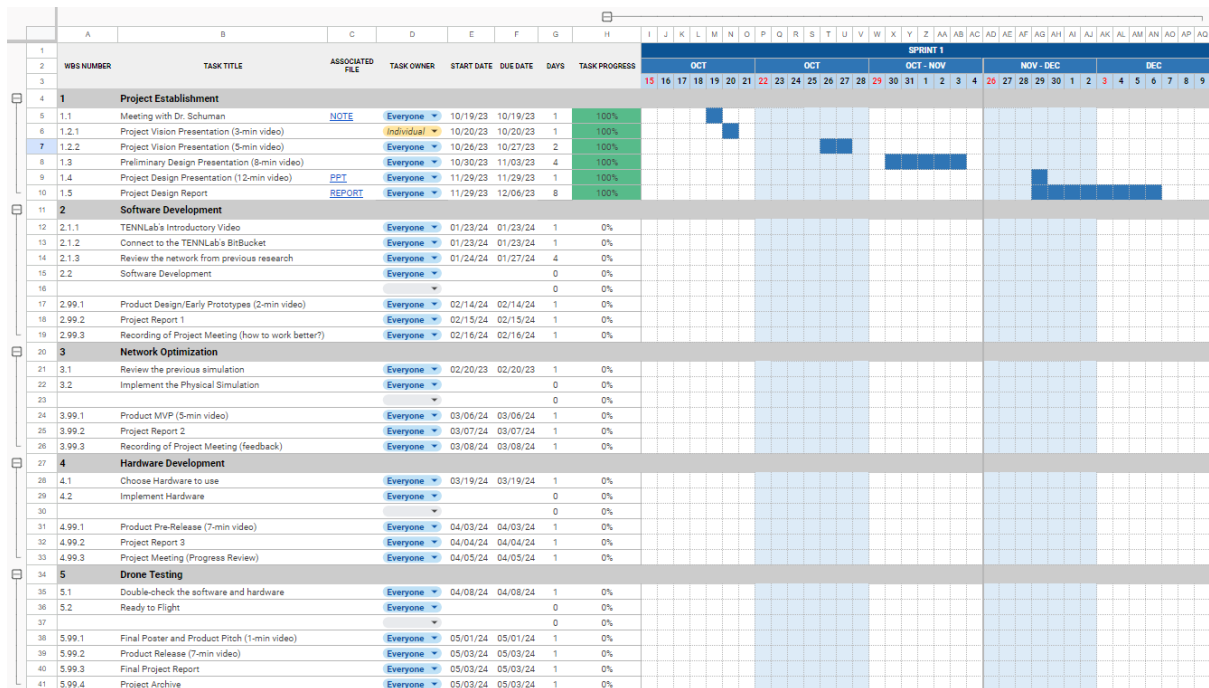


Figure 5: Part of the Gantt Table of the Project

The Gantt chart has been modified to improve accessibility and make schedule management more convenient by modifying the template provided by Google Spreadsheet to automatically fill the table when someone enters the start and due dates. It also has the advantage of being able to easily determine who is working on what task. We expect that by updating the Gantt chart, we will be able to easily check the progress of the project and through this, we will be able to control the progress of the project.

BUDGET

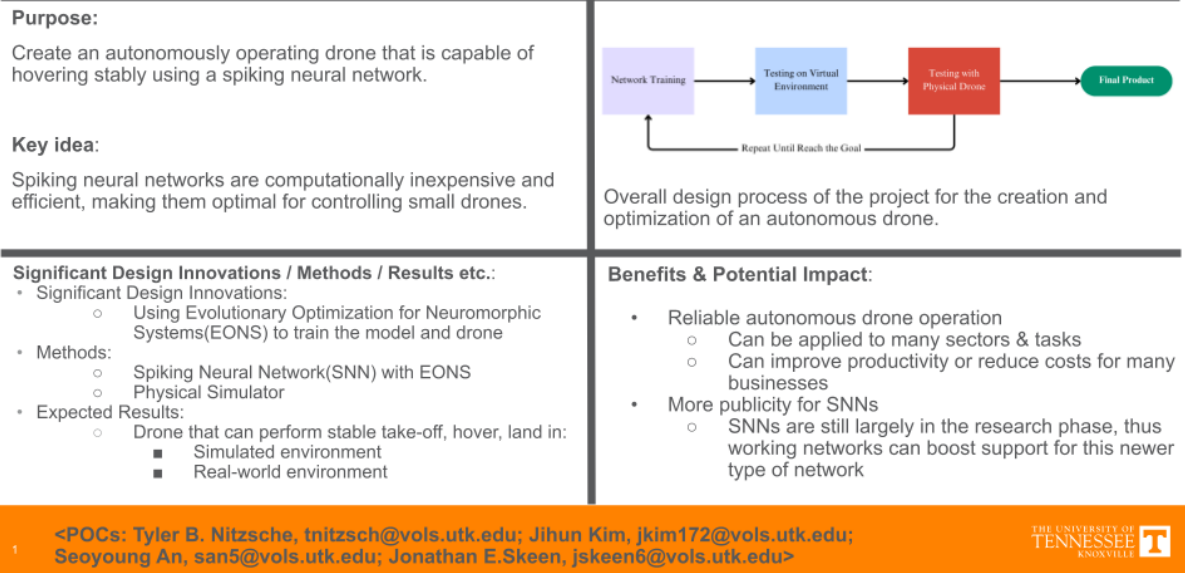
We don't have a definitive budget. To order necessary parts or equipment, we will create a list of parts, their prices, and associated links that lead to where the parts can be bought. The list will be confirmed by our advisor, Dr. Schuman. Then, we will send an email to eecs-orders@utk.edu and CC the teaching assistant of the COSC 401 course, including the name of our project, team lead, list of items, price range, and explanation of why we need them. The possible items that we might need to buy are another drone and sensors like gyroscopes for the drone. Once we start working with the physical drone, we will be able to figure out whether we need to order additional parts or not.

Reference

- [1] Bitcraze. Crazyflie 2.1.
- [2] Catherine D Schuman, Shruti R Kulkarni, Maryam Parsa, J Parker Mitchell, Prasanna Date, and Bill Kay. Opportunities for neuromorphic computing algorithms and applications. *Nature Computational Science*, 2(1):10–19, 2022.
- [3] C. D. Schuman, J. P. Mitchell, R. M. Patton, T. E. Potok, and J. S. Plank. Evolutionary optimization for neuromorphic systems. In *NICE: Neuro-Inspired Computational Elements Workshop*, 2020.
- [4] Simon D. Levy. Multicoptersim, 2023.
- [5] Simon D. Levy. gym-copter, 2023.

APPENDIX

< Spiking Neural Network Drones >



<POCs: Tyler B. Nitzsche, tnitzsch@voils.utk.edu; Jihun Kim, jkim172@voils.utk.edu; Seoyoung An, san5@voils.utk.edu; Jonathan E.Skeen, jskeen6@voils.utk.edu>



Figure 6: Project Quad Chart

Table 1: Milestone in Each Sprint

Sprint 1	10/02/2023 – 01/22/2024
<p>Sprint 1 is the longest, running from the start of the project until the start of Spring 2024. During Sprint 1, group members identify the size of the project, set project goals, and devise detailed goals for this goal.</p> <p>During this period, group members met with our mentor, Dr. Schuman, to hear about the project and gain insight into the project. This information includes software resources, hardware resources, place support, and funding available to us.</p>	
Sprint 2	01/21/2024 – 02/17/2024
<p>Sprint 2 is the period from the start of the 2024 spring semester to the filming of the first video. The goal of Sprint 2 is to produce an early prototype. During this period, group members can increase their understanding of TENNLAB's resources, including EONS, through TENNLAB's introductory videos, and acquire basic knowledge to implement the software needed for the project. After this, we will find out the level of the network that has already been implemented, establish directions on how to improve it, and begin software development.</p>	
Sprint 3	02/18/2024 – 03/08/2024
<p>Sprint 3 is the period from the first video recording to the second video recording. The goal of Sprint 3 is the complete (or near complete) implementation of the software for hardware. During this period, we will seek to improve the early prototype implemented in Sprint 2 through network optimization and training. In this process, we will utilize a physics simulator. We can either use the simulator designed by the previous senior design course or one that has been developed by Simon Levy.</p>	
Sprint 4	03/17/2024 – 04/06/2024
<p>Sprint 4 is the period from the second video recording to the third video recording. During this period, we will work on adapting the software to the hardware. We will select a small Crazyflie drone suitable for our project and apply our software to it. These processes may include hardware design as well as hardware implementation. The goal of Sprint 4 is the successful implementation of the hardware, i.e., the implementation of a viable drone.</p>	
Sprint 5	04/07/2024 – 05/03/2024
<p>Sprint 5 is the period from the third video recording to the end of the project. The goal of Sprint 5 is to prepare the drone for flight and presentation. Sprint 5 includes the final tuning of the hardware and commissioning of the actual machine. We will evaluate our work by flying our drone. Considering that the flights of these drones require permits, we will do our best to ensure that these flights are successful from takeoff to landing.</p>	